

Reviewer-defense moat for assay-ready translation

The model was optimized on current labels.

Top-13 label-aware result is framed as a decision-support operating point; source-heldout/external validation is separated.

MD cannot prove immunogenicity.

MD is explicitly used only for structural plausibility and control selection; activation assays are required.

TCR-aware prediction is not general because most datasets lack paired TCR.

TCR is a case-study/expert branch; main pMHC ranking is kept separate.

WT cross-reactivity was not tested.

WT controls are now explicit: KRAS WT GAGGVGKSAL and TP53 WT HMTEVVRRC; no immunogenicity claim before WT/decoy gates

These are literature-known antigens, not novel discovery.

Use them as anchor case studies to validate the pipeline's triage logic, while keeping discovery claims for independent candidate